



JPMC #2: Self-Hosted End-to-End Question Answering Bot AI Studio Final Presentation

Break Through Tech New York @Cornell Tech
[12.13.2023]



Introductions



Meet Our Team!




Alyssa Choi
Stony Brook University



Laurretta Martin
Rutgers University



Sonia Tan
Stony Brook University



Hanqi Lin
CUNY Queens College



Fayha Farooqi
Stevens Institute of Technology

Our AI Studio TA and Challenge Advisors



Sakshi Pandey
AI Studio TA



Adedapo Alabi
Challenge Advisor



Presentation Agenda

1. **AI Studio Project Overview**
2. **Architecture**
3. **Implementation**
4. **Vector Search LLM**
5. **Final Implementation and Demonstration**




AI Studio Project Overview



Self-Hosted Product-Centric Project

Utilizing a pre-trained LLM of public data

End-to-end implementation





Our Goal

1. Design an **end-to-end self-hosted application** that allows users to retrieve information from a private pdf document
2. Redesign to implement **multiple documents** to be inputted into the chatbot
3. Allow JPMC to deliver recommendations and information to clients in-house without third-parties involved
 - a. Companies with financial data don't have the liberty to send data to OpenAI due to the necessity of **data security** regarding financial statements
 - b. Federal Trade Commission - Financial Privacy Regulations



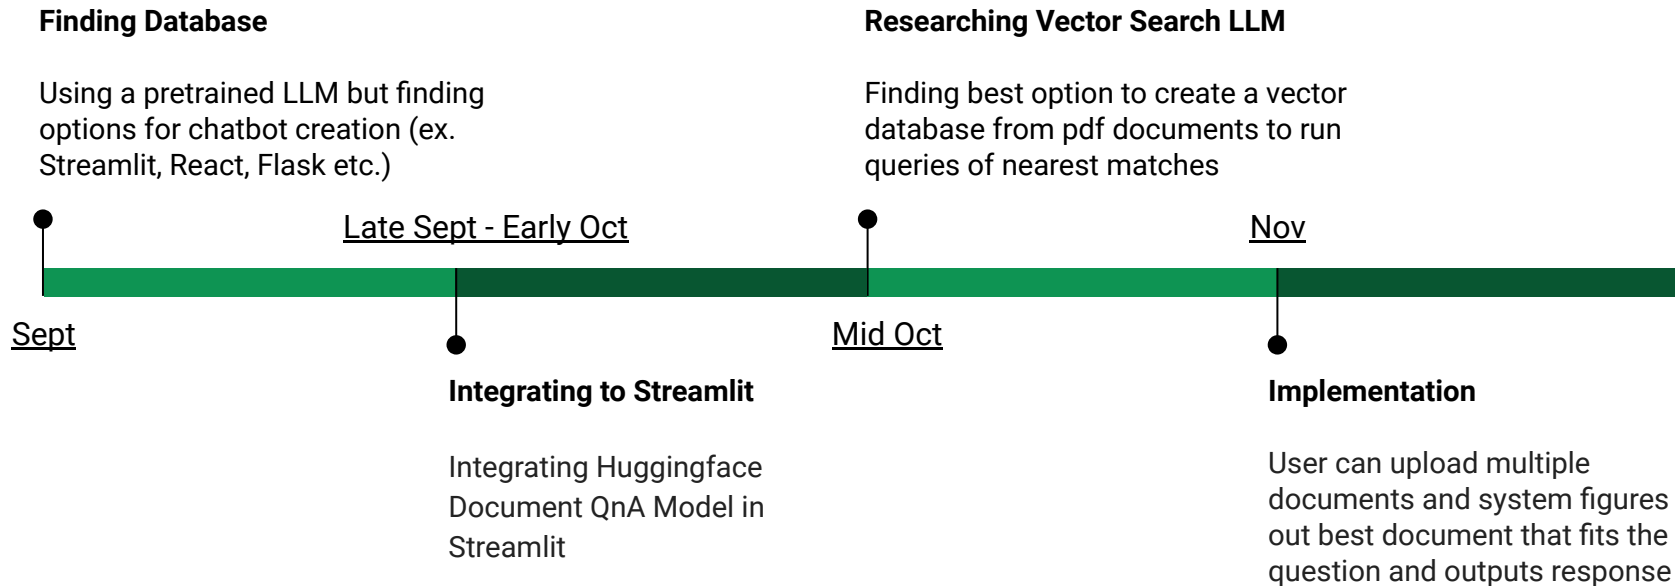
Business Impact

- Worked with JP Morgan Chase on this project
- With explosions of LLMs such as ChatGPT, chatbots are now able to have human-like conversations
- Companies like JPMC would benefit from chatbots that allow users to talk to these documents without sharing any data with external services



Our Approach

Decided on the following approach based on Challenge Advisor guidance. Before starting the project, we researched different architectures and frameworks that are feasible for us to implement.






How can financial companies abide by data regulations while still utilizing AI/ML/DL to make processes more efficient?

- Create data processing programs in-house!
 - Maintains **data integrity** and safety
 - Abides by financial institution data privacy laws
 - Gives firms the **opportunity to cater their created software product** to their own needs
 - **Streamlined** data collection and processing



Resources We Leveraged

- HuggingFace
- Streamlit
- Github

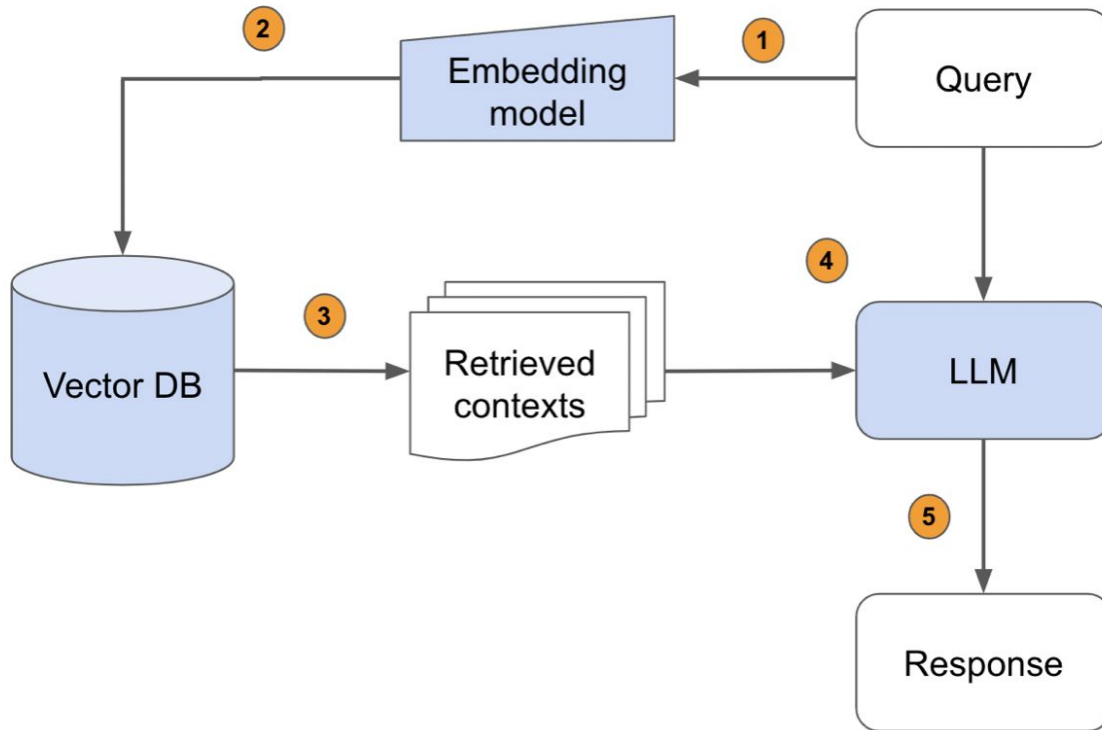




Architecture



Program Architecture





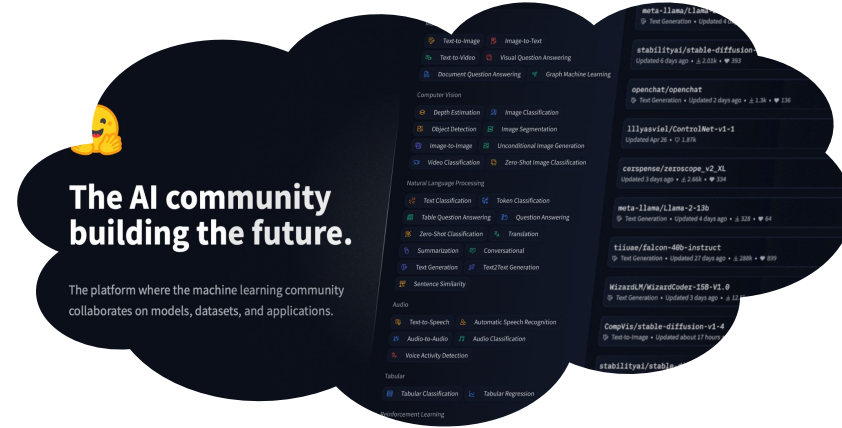
Implementation

Model Research



- **HuggingFace**

- Building and deploying machine learning models
- Includes open-source pre-trained models
- Ideal to test chatbot on in regards to basic information
- Involves a wide range of datasets to test different types of questions on





Streamlit is an open-source app
framework for machine learning and data
science web apps



Streamlit



DocQuery Implementation

- Takes a single image file or single page PDF
- Takes a question from the user
- Returns an answer using DocQuery

Choose a file



Drag and drop file here

Limit 200MB per file

Browse files



wordpress-pdf-invoice-plugin-sample.pdf 42.6KB



How may I help you?



What is the invoice number?



INV-3337

Your message





RoBERTa Implementation

- **RoBERTa Overview**

- Self-supervised transformers model pretrained on large English texts
- Bidirectional representation learning

- **Implementation**

- Takes a single PDF file and extracts the text from the file
- Takes a question from the user
- Returns an answer using the RoBERTa model
- Demo



```
# File uploader widget
uploaded_file = st.file_uploader("Choose a file")

# Upload file to temp folder and get file path
if uploaded_file:
    temp_dir = tempfile.mkdtemp()
    path = os.path.join(temp_dir, uploaded_file.name)
    with open(path, "wb") as f:
        f.write(uploaded_file.getvalue())

# Store LLM generated responses
if "messages" not in st.session_state.keys():
    st.session_state.messages = [{"role": "assistant", "content": "How may I help you?"}]

# Display chat messages
for message in st.session_state.messages:
    with st.chat_message(message["role"]):
        st.write(message["content"])

# Function for generating response
def generate_response(path, prompt_input):
    p = pipeline("document-question-answering")
    doc = document.load_document(path)
    return p(question=prompt_input, **doc.context)[0]["answer"]

# User-provided prompt
if prompt := st.chat_input():
    st.session_state.messages.append({"role": "user", "content": prompt})
    with st.chat_message("user"):
        st.write(prompt)

# Generate a new response if last message is not from assistant
if st.session_state.messages[-1]["role"] != "assistant":
    with st.chat_message("assistant"):
        with st.spinner("Thinking..."):
            response = generate_response(path, prompt)
        st.write(response)
    message = {"role": "assistant", "content": response}
    st.session_state.messages.append(message)
```

```
import streamlit as st
from hugchat import hugchat
from transformers import pipeline
from pypdf import PdfReader
import io

st.set_page_config(page_title="AI Finance Chatbot🤖")
st.header('AI Finance Chatbot🤖')

uploaded_file = st.file_uploader('Choose your .pdf file', type="pdf")
if uploaded_file is not None:
    file_contents = uploaded_file.read()
    remote_file_bytes = io.BytesIO(file_contents)
    pdfdoc_remote = PdfReader(remote_file_bytes)

    pdf_text = ""

    for i in range(len(pdfdoc_remote.pages)):
        print(i)
        page = pdfdoc_remote.pages[i]
        page_content = page.extract_text()
        pdf_text += page_content

    print(pdf_text)

    nlp = pipeline(
        "question-answering",
        model="deepset/roberta-base-squad2",
        tokenizer="deepset/roberta-base-squad2",
    )

    if "messages" not in st.session_state.keys():
        st.session_state.messages = [
            {"role": "assistant", "content": "Ask me a question about the document!"}
        ]
```



Vector Search LLMs



Vector Search in Large Language Models (LLMs): A Comprehensive Overview

- **What is Vector Search?**

- Turns text into numerical vectors representing its meaning.
- Helps find relevant answers in a huge amount of data.

- **How Does it Work?**

- Input Conversion: Changes your question into a vector.
- Matching: Compares your vector with stored vectors to find the best match.

- **Benefits**

- Fast: Quickly sifts through large data.
- Accurate: Understands context, not just keywords.



Implementing Vector Search in Chatbots

Input as Query: AI interprets request as a query. Question is now converted into a numerical vector.

Semantic Understanding: LLM analyzes the text's content, breaking down into simpler concepts.

Response Generation: The AI then generates a response, akin to how vector search matches a query vector with the most relevant document vectors, the essence of the original text is preserved.

student input: “ to make my HW more “readable and to the point”

gpt respond: In order to make your HW more “concise and reader-friendly”.....



Final Implementation and Demonstration



Txt-AI Implementation

- Converts pdf files into txt
- Creates vector embeddings of all files and stores in a vector database
- Searches database for file with closest context

Choose your files



Drag and drop files here

Limit 200MB per file

Browse files



Reference Letter Template_0.pdf 102.9KB



wordpress-pdf-invoice-plugin-sample.pdf 42.6KB



How may I help you?



What is the invoice number?



INV-3337

Your message





References and Resources

- [impira/layoutlm-document-qa · Hugging Face](#)
- [GitHub - impira/docquery: An easy way to extract information from documents](#)
- [deepset/roberta-base-squad2 · Hugging Face](#)
- [GitHub - neuml/txtai: 💡 All-in-one open-source embeddings database for semantic search, LLM orchestration and language model workflows](#)
- [How to build an LLM-powered ChatBot with Streamlit](#)
- [GitHub - py-pdf/pypdf: A pure-python PDF library capable of splitting, merging, cropping, and transforming the pages of PDF files](#)



Thank You!

**Please advise of any questions
or comments, we'd love to
hear your feedback!**